
Counterfactual Audits Must Precede Conformal Trust for Multimodal Medical Claim Verifiers

Anonymous Authors¹

Abstract

Multimodal claim verifiers are placed downstream of chest-X-ray report generators as a safety layer, and several recent verifiers report above 90% accuracy on hallucination-detection benchmarks. We show that those numbers are attainable by a verifier that never uses the image. We call this failure mode evidence-blindness. We define three counterfactual perturbations of a verifier’s input and measure the accuracy drop: image zeroing (IMG), within-batch evidence shuffling (ESG), and left–right image flipping on laterality claims (IPG). On a public-data benchmark, the configuration with the highest validation accuracy is evidence-blind on IMG (2.17 pp); adding an image-masked consistency loss and an adversarial hypothesis-only filter raises IMG to 69.24 pp at matched test accuracy. The mitigation is brittle: under leave-one-site-out training, IMG regresses to baseline on the held-out site. We further wrap inverted conformal Benjamini–Hochberg selection around each verifier and find that every evidence-blind configuration returns an empty trusted set; only the non-blind verifier yields a populated one. A conformal guarantee on an evidence-blind verifier has nothing to certify. Counterfactual sensitivity testing should precede both aggregate-accuracy reporting and conformal trust wrapping for any multimodal medical verifier.

1. Introduction

Radiology report generators are entering radiologist review queues as draft producers, with reported hallucination rates of 8–15% (Zhang et al., 2025; Hardy et al., 2025). The community response has been to insert a claim-level verifier downstream, with several systems reporting AUROC above 90% on aggregate hallucination-detection benchmarks. Discussion has moved toward which conformal guarantee should wrap these verifiers (Mohri & Hashimoto, 2024; Angelopoulos & Bates, 2021; Gui et al., 2024; Li et al., 2025).

Aggregate accuracy does not distinguish a verifier that uses the image from one that exploits shortcut structure in the claim and evidence text. Gururangan et al. (2018) and Poliak et al. (2018) established this for natural language inference: hypothesis-only classifiers approached the accuracy of full-input classifiers on SNLI and MultiNLI without observing the premise, which reshaped NLI evaluation practice. Restrepo et al. (2025) address the radiology instance directly with Selective Modality Shifting on MIMIC-CXR and FairVLMed, and report textual dependency across four generalist VLMs and two medical-specialist VLMs; that work explicitly leaves training-time mitigation to future work, and the conformal-selection literature wraps an existing predictor without modifying it.

This paper makes two contributions. (i) We define evidence-blindness through three counterfactual metrics — IMG, ESG, and a laterality-flip variant IPG that does not appear in the SMS family — that apply to any binary multimodal verifier including closed-weight APIs. (ii) We train a cross-modal verifier on a public, non-credentialed chest-radiograph benchmark; its highest-validation-accuracy configuration is evidence-blind on IMG. A consistency loss combined with an adversarial hypothesis-only filter raises IMG from 2.17 to 69.24 pp at matched accuracy. An inverted conformal Benjamini–Hochberg selection procedure collapses to an empty trusted set on every evidence-blind configuration we tested, and returns a populated trusted set only on the non-blind one. The counterfactual diagnostic should run before conformal

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

wrapping in any deployment pipeline.

2. Methods and Benchmark

Counterfactual diagnostics. Let f be a binary multimodal verifier mapping $(x_{\text{img}}, x_{\text{claim}}, x_{\text{evid}})$ to $\{\text{sup}, \text{con}\}$. On a labelled evaluation set $\mathcal{D}$:

$$\text{IMG}(f) = \text{acc}(f, \mathcal{D}) - \text{acc}(f, \mathcal{D}^{\text{img}=0}), \quad (1)$$

$$\text{ESG}(f) = \text{acc}(f, \mathcal{D}) - \mathbb{E}_{\pi}[\text{acc}(f, \mathcal{D}^{\text{evid}=\pi})], \quad (2)$$

$$\text{IPG}(f) = \text{acc}(f, \mathcal{D}_{\text{lat}}) - \text{acc}(f, \mathcal{D}_{\text{lat}}^{\text{flip}}), \quad (3)$$

in percentage points. $\mathcal{D}^{\text{img}=0}$ replaces each image with a zero tensor (we use literal zeros for the IMG diagnostic since the verifier inputs are normalised CLIP features, not pixel intensities); $\mathcal{D}^{\text{evid}=\pi}$ replaces evidence text with $x_{\text{evid}}^{(\pi(i))}$ under a within-batch derangement π averaged over three seeds; $\mathcal{D}_{\text{lat}}^{\text{flip}}$ horizontally flips the image on the subset of claims matching `\b(left|right|bilateral)\b`. We label f evidence-blind if $\text{IMG}(f) < 5$ pp or $\text{ESG}(f) < 5$ pp; threshold sensitivity over $\{3, 5, 7, 10\}$ pp is reported in the supplementary appendix. The diagnostic requires only input substitution and applies to closed-weight verifiers exposed via API.

Cross-modal verifier. Images at 224×224 pass through BiomedCLIP ViT-B/16 (Zhang et al., 2024); the top four of twelve transformer blocks are trainable. Claim and evidence tokens are jointly encoded by RoBERTa-large (Liu et al., 2019) via the standard text-pair API; the top eight of twenty-four layers are trainable. Both modalities project to $d = 768$. A learnable [VERDICT] token is prepended at position zero of the fused sequence, which passes through four bidirectional cross-modal transformer layers with twelve heads each. Three heads attach to the verdict-token output: a binary verdict head (2-way softmax), a support-score head (sigmoid scalar, used downstream for conformal calibration), and a 14×14 per-patch grounding head (sigmoid map). The model has ≈ 480 M total parameters of which 122 M are trainable. Full optimisation details and the loss equation are in the supplementary appendix.

Composite loss and hypothesis-only filter. The training loss is a five-term composite: classification cross-entropy ($\mathcal{L}_{\text{cls}}$); bounding-box-supervised patch grounding ($\mathcal{L}_{\text{ground}}$, $\lambda = 0.5$); an image-masked consistency margin $\mathcal{L}_{\text{consist}} = \text{ReLU}(0.2 - (p_{\text{full}}(y) - p_{\text{img}=0}(y)))$ ($\lambda = 0.3$); a contrastive evidence margin restricted to sup anchors ($\lambda = 0.3$); and an MC-dropout uncertainty regulariser ($\lambda = 0.1$). The adversarial hypothesis-only (HO) filter is a separate

Table 1. Counterfactual diagnostic. Top: in-distribution test split. Bottom: cross-site leave-one-out (the v5.3 loss composition retrained for two epochs on the non-held-out sites and evaluated on the held-out site). A configuration is evidence-blind if $\text{IMG} < 5$ pp or $\text{ESG} < 5$ pp.

Config	Test acc	IMG	ESG	Blind?
v5.0-base	0.909	2.03	19.71	yes
v5.1-ground	0.921	2.17	21.13	yes
v5.2-real	0.923	62.90	18.74	no
v5.3-contrast	0.925	69.24	18.56	no
v6.0-retrain	0.911	70.21	20.01	no
Cross-site leave-one-out (held-out site):				
held out: OpenI	0.853	3.21	-0.02	yes
held out: ChestX-Det10	0.660	28.41	2.69	yes

RoBERTa-large cross-encoder trained for one epoch on $(\text{claim}, \text{evidence}) \rightarrow y$ pairs with no image input. At training time, rows on which the HO filter’s gold-class probability exceeds 0.7 have $\mathcal{L}_{\text{cls}}$ downweighted by $0.2\times$; the other four terms retain full weight. Loss-drop ablations (supp. appendix) show $\mathcal{L}_{\text{consist}}$ and the HO filter are the two load-bearing terms: removing either collapses IMG to ≤ 6 pp.

Benchmark. GroundBench-v6 combines three public sources: OpenI (Demner-Fushman et al., 2016) ($\sim 4,000$ frontal radiographs), ChestX-Det10 (Liu et al., 2020) ($\sim 3,500$ images with pixel annotations), and PadChest-GR (Castro et al., 2024) (4,555 bilingual studies with per-sentence radiologist boxes). Patient-stratified splits at 70/10/10/10 yield 28,361 train, 4,046 validation, 4,046 calibration, and 4,132 test claims. For real-RRG evaluation, 1,500 findings reports from MAIRA-2 (Bannur et al., 2024), CheXagent-2 (Chen et al., 2024), and MedGemma-4B-IT (Sellergren et al., 2025) are decomposed into 2,707 atomic claims and silver-labelled by a three-LLM judge ensemble (Ostmeier et al., 2024). Full benchmark detail is in the supplementary appendix.

3. Results

The classification baseline is evidence-blind. Table 1 reports the diagnostic on the 2,906-claim in-distribution test split. The cross-entropy baseline v5.0 reaches 90.9% test accuracy at $\text{IMG} = 2.03$ pp; v5.1 adds patch grounding, reaches the highest validation accuracy (91.1%), and stays evidence-blind at $\text{IMG} = 2.17$ pp. A practitioner selecting on aggregate accuracy alone would ship the most blind of the configurations we trained.

Consistency + HO filter is the load-bearing intervention. v5.2 adds both. IMG rises from 2.17 to 62.90 pp;




Figure 1. Aggregate accuracy and image-grounding are independent axes. Circles: trained configurations. Squares: zero-shot detectors on a 500-claim real-RRG subset. Shaded band: evidence-blind zone (IMG < 5 pp). The highest-validation-accuracy configurations (v5.0, v5.1) sit in the blind zone at 90–92% accuracy; zero-shot baselines BiomedCLIP and MedGemma-4B-IT are blind on the ESG axis at intermediate accuracy.

the zero-image accuracy collapses to 29.4% (below random) because the consistency loss rewards verdict divergence under image masking, while aggregate accuracy rises to 92.3%. Adding contrastive evidence (v5.3) lifts IMG to 69.24 pp at 92.5% accuracy. A confirmatory v6.0 retrain reaches 70.21 pp.

Field audit on real-RRG hallucinations. We extend the diagnostic to 500 silver-labelled real-RRG claims with four zero-shot detectors (Fig. 1; full numbers in the supplementary appendix). BiomedCLIP zero-shot reaches 0.538 accuracy with IMG = 8.0 pp and ESG = 0.0 pp (blind on ESG). MedGemma-4B-IT used as a verifier reaches 0.758 with IMG = 14.0 pp and ESG = 2.0 pp (blind on ESG). MAIRA-2 used zero-shot clears both thresholds (IMG = 32.2 pp, ESG = 6.6 pp) but at near-chance accuracy (0.498). The combination matters: a verifier can fail one axis without failing the other, and aggregate accuracy alone does not predict the joint outcome.

Cross-site regression and deployment alarm. Two leave-one-site-out configurations trained under the v5.3 loss composition show the mitigation does not transfer (Table 1, bottom). The held-out OpenI run regresses to IMG = 3.21 pp (within 1.2 pp of the v5.0 baseline); held-out ChestX-Det10 keeps most of the IMG mitigation but drops to ESG = 2.69 pp, below threshold. The same three metrics that flag training-time blindness flag post-deployment blindness. This is the second role of the same diagnostic.

Conformal selection collapses on blind verifiers. We wrap an inverted conformal Benjamini–Hochberg procedure around each verifier’s support score, calibrated on a held-out 4,046-claim calibration split at target $\alpha = 0.10$. Only v5.3-contrast produces a populated trusted set: $n_{\text{green}} = 985$ claims with achieved FDR = 0.91% and power 39.6%. Configurations v5.0, v5.1, and v5.2 collapse to $n_{\text{green}} = 0$: v5.0 and v5.1 because they are evidence-blind, and v5.2 because the joint distribution of (support score, true label) is not separable enough below the target level (supp. appendix). A conformal guarantee on a blind verifier has nothing to certify.

4. Discussion

The cross-entropy baseline reports a competitive 90.9% test accuracy while being image-blind at IMG = 2.03 pp. The consistency loss and HO filter act on disjoint failure modes: the consistency loss penalises verdict-invariance under image masking, and the HO filter downweights training rows the text-only model already solves. Patch grounding alone (v5.1) does not move IMG, mirroring NLI partial-input findings (Gururangan et al., 2018; Poliak et al., 2018) that auxiliary supervision on input attribution does not by itself prevent shortcut exploitation.

The mitigation indexes on training-distribution shortcuts. Under $\text{site}_{\text{train}} \neq \text{site}_{\text{test}}$, the shortcut structure changes and the mitigation has nothing to suppress — but the same diagnostic catches the post-shift blindness. This makes the counterfactual diagnostic the appropriate alarm signal in both training-time selection and deployment-time monitoring; the appropriate fix for the cross-site regime is distribution-shift-aware methodology (e.g., weighted conformal inference under covariate shift (Tibshirani et al., 2019)) layered on top of the same diagnostic.

The conformal-selection collapse on blind verifiers sharpens the Mohri–Hashimoto conformal-factuality guarantee (Mohri & Hashimoto, 2024): marginal coverage is preserved on whatever signal the underlying verifier emits, but if that signal is image-blind the resulting trusted set has no clinical content. Wrapping a frozen verifier in conformal selection is necessary but not sufficient for trustworthy deployment.

5. Limitations and Conclusion

Three limitations bound our claims. (L1) Scale. The benchmark spans approximately 7,500 images and 25,176 training claims — enough to expose evidence-blindness as a measurable failure mode but not enough

to estimate its prevalence at population scale. (L2) Single-seed training. All training runs use a single random seed; multi-seed replication is deferred. (L3) Spatial residual. Laterality-turning claims are insensitive to horizontal flips ($\text{IPG} \in [-0.88, +0.88]$ pp across all configurations); the consistency loss targets image presence rather than spatial arrangement, and closing this residual would require an image-encoder-level intervention.

We introduced a counterfactual diagnostic for multimodal medical claim verification (IMG, ESG, IPG), a training-time recipe that closes the in-distribution IMG from 2.03 pp to 69.24 pp at +1.6 pp accuracy, and the empirical observation that inverted conformal selection collapses to an empty trusted set on every evidence-blind verifier we tested. Counterfactual evidence-sensitivity testing should precede aggregate-accuracy reporting and conformal trust-selection for any multimodal medical verifier. Benchmark maintainers and submission checklists in this area should report partial-input baselines alongside aggregate accuracy — as the NLI community did after Gururangan et al. (2018) and Poliak et al. (2018).

References

- Angelopoulos, A. N. and Bates, S. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. arXiv preprint arXiv:2107.07511, 2021. URL <https://arxiv.org/abs/2107.07511>.
- Bannur, S., Bouzid, K., Castro, D. C., Schwaighofer, A., Thieme, A., Bond-Taylor, S., Ilse, M., Pérez-García, F., Salvatelli, V., Sharma, H., Meissen, F., Ranjit, M., Srivastav, S., Gong, J., Codella, N. C. F., Falck, F., Oktay, O., Lungren, M. P., Wetscherek, M. T., Alvarez-Valle, J., and Hyland, S. L. MAIRA-2: Grounded radiology report generation. arXiv preprint arXiv:2406.04449, 2024. URL <https://arxiv.org/abs/2406.04449>. Microsoft Research; research-only licence.
- Castro, D. C., Bustos, A., Bannur, S., Hyland, S. L., Bouzid, K., Wetscherek, M. T., Sánchez-Valverde, M. D., Jaques-Pérez, L., Pérez-Rodríguez, L., Takeda, K., Salinas, J., Alvarez-Valle, J., Galant Herrero, J., and Pertusa, A. PadChest-GR: A bilingual chest X-Ray dataset for grounded radiology report generation. arXiv preprint arXiv:2411.05085, 2024. URL <https://arxiv.org/abs/2411.05085>. Also published in NEJM AI (2025).
- Chen, Z., Varma, M., Xu, J., Paschali, M., Van Veen, D., Johnston, A., Youssef, A., Blankemeier, L., Bluethgen, C., Altmayer, S., et al. A vision-language foundation model to enhance efficiency of chest X-ray interpretation. arXiv preprint arXiv:2401.12208, 2024. URL <https://arxiv.org/abs/2401.12208>. Stanford-AIMI. The model is referred to as CheXagent/CheXagent-2 in the system documentation.
- Demner-Fushman, D., Kohli, M. D., Rosenman, M. B., Shooshan, S. E., Rodriguez, L., Antani, S., Thoma, G. R., and McDonald, C. J. Preparing a collection of radiology examinations for distribution and retrieval. Journal of the American Medical Informatics Association, 23(2):304–310, 2016. URL <https://openi.nlm.nih.gov/>. OpenI Chest X-ray collection (Indiana University).
- Gui, Y., Jin, Y., and Ren, Z. Conformal alignment: Knowing when to trust foundation models with guarantees. arXiv preprint arXiv:2405.10301, 2024. URL <https://arxiv.org/abs/2405.10301>. Reference-data-trained alignment predictor for FDR-controlled output selection. Applied to QA + radiology RG. Operates on a frozen predictor.
- Gururangan, S., Swayamdipta, S., Levy, O., Schwartz, R., Bowman, S., and Smith, N. A. Annotation artifacts in natural language inference data. In NAACL-HLT, 2018. URL <https://aclanthology.org/N18-2017/>.
- Hardy, R., Kim, S. E., Ro, D. H., and Rajpurkar, P. ReXTrust: A model for fine-grained hallucination detection in AI-generated radiology reports. In Proceedings of the First AAAI Bridge Program on AI for Medicine and Healthcare, volume 281 of Proceedings of Machine Learning Research, 2025. URL <https://proceedings.mlr.press/v281/hardy25a.html>. White-box hidden-state probe for hallucination detection.
- Li, Z., Yan, C., Jackson, N. J., Cui, W., Li, B., Zhang, J., and Malin, B. A. Towards statistical factuality guarantee for large vision-language models. arXiv preprint arXiv:2502.20560, 2025. URL <https://arxiv.org/abs/2502.20560>. Conformal claim filtering on LLaVa-1.5 (frozen LVLm).
- Liu, J., Lian, J., and Yu, Y. ChestX-Det10: Chest X-ray dataset on detection of thoracic abnormalities. arXiv preprint arXiv:2006.10550, 2020. arXiv-only release. Bounding-box annotations; no paired findings text.
- Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer, L.,

- and Stoyanov, V. RoBERTa: A robustly optimized BERT pretraining approach. arXiv preprint arXiv:1907.11692, 2019. URL <https://arxiv.org/abs/1907.11692>.
- Mohri, C. and Hashimoto, T. Language models with conformal factuality guarantees. In International Conference on Machine Learning (ICML), 2024. URL <https://arxiv.org/abs/2402.10978>.
- Ostmeier, S., Xu, J., Chen, Z., Varma, M., Blanke-meier, L., Bluethgen, C., Michalson, A. E., Moseley, M., Langlotz, C., Chaudhari, A. S., and Delbrouck, J.-B. GREEN: Generative radiology report evaluation and error notation. In Findings of the Association for Computational Linguistics: EMNLP 2024, 2024. URL <https://arxiv.org/abs/2405.03595>.
- Poliak, A., Naradowsky, J., Haldar, A., Rudinger, R., and Van Durme, B. Hypothesis only baselines in natural language inference. In *SEM, 2018. URL <https://aclanthology.org/S18-2023/>.
- Restrepo, D., Ktena, I., Vakalopoulou, M., Christodoulidis, S., and Ferrante, E. On the risk of misleading reports: Diagnosing textual biases in multimodal clinical AI. arXiv preprint arXiv:2508.00171, 2025. URL <https://arxiv.org/abs/2508.00171>. Accepted at MICCAI 2025 1st Workshop on Multimodal LLMs in Clinical Practice.
- Sellergren, A., Kazemzadeh, S., Jaroensri, T., Kiraly, A., Traverse, M., Kohlberger, T., Xu, S., Jamil, F., Hughes, C., Lau, C., et al. MedGemma technical report. arXiv preprint arXiv:2507.05201, 2025. URL <https://arxiv.org/abs/2507.05201>. Google Health AI Developer Foundations; model weights at <https://huggingface.co/google/medgemma-4b-it> under HAI-DEF research-only licence.
- Tibshirani, R. J., Foygel Barber, R., Candès, E. J., and Ramdas, A. Conformal prediction under covariate shift. In Advances in Neural Information Processing Systems (NeurIPS), 2019. URL <https://arxiv.org/abs/1904.06019>.
- Zhang, S., Xu, Y., Usuyama, N., Xu, H., Bagga, J., et al. BiomedCLIP: a multimodal biomedical foundation model pretrained from fifteen million scientific image-text pairs. arXiv:2303.00915 / Nature Communications, 2024. URL <https://arxiv.org/abs/2303.00915>. Biomedical image encoder.
- Zhang, S., Sambara, S., Banerjee, O., Acosta, J., Fahrner, L. J., and Rajpurkar, P. RadFlag: A black-box hallucination detection method for medical vision-language models. In Proceedings of the 4th Machine Learning for Health Symposium, PMLR 259:1087-1103, 2025. URL <https://arxiv.org/abs/2411.00299>. Self-agreement-based hallucination detection.